


## Sec B Team 8

Kaggle Team name: Sec B Team 8


Best LR model: lr\_default, accuracy = 0.880 , Best RF model: rf\_default, accuracy = 0.841

### I. Default features

Best for LR: `cs = np.arange(0.1, 3.5, 0.1)` , accuracy = 0.880



Best for RF: `n_estimators = np.arange(20,100,20)` , accuracy = 0.841



First, to optimize the hyper-parameters of logistic regression, I modified `get_tuned_lr` to several ranges, and compared the plots they generated. The final range I picked was `cs = np.arange(0.1, 3.5, 0.1)`, because the accuracy is high, and also it offers the smallest difference between the training and testing accuracy, so we can better generalize our model for unseen data.

Second, to optimize the parameters for random forest, I also modified `get_tuned_rf`. `N_estimators` for several times, and I found as the number increased, the accuracy also increased, so the optimal range with highest accuracy I found was actually (50, 500, 50), all of its accuracy scores are 1 for the train dataset. But bigger numbers take a very long time to run, so I start decreasing the numbers and compare the plots, then the final range I got was (20,100,20) .

### II. Custom features

lr\_custom: accuracy = 0.823, rf\_custom: accuracy = 0.837

I used 'pos\_count': self.pos\_count(review), and 'neg\_count': self.neg\_count(review) as one group, then tried different combinations with others like word\_count. Then I tried CountVectorizer() with different ngram\_ranges, and compared the results to TfidfVectorizer() with different ngram\_ranges also. I found there wasn't a big difference between CountVectorizer() and TfidfVectorizer(), and the accuracy increases only a little bit when changing the ngrams, but runs very slow, so I just used CountVectorizer() with default ngram ranges.

### **III. Analysis**

Classifier1 Metrics

Accuracy: 0.9546

Precision: 0.9534753810549836

Recall: 0.95584

F1: 0.9546562262794135

Classifier2 Metrics

Accuracy: 0.96924

Precision: 0.9674822666772934

Recall: 0.97112

F1: 0.9692977202858625

Classifier3 Metrics

Accuracy: 0.83808

Precision: 0.8276476973174135

Recall: 0.854

F1: 0.8406173714465706

Classifier4 Metrics

Accuracy: 0.96872

Precision: 0.9683453237410072

Recall: 0.96912

F1: 0.968732506997201

The best performance classifier is Classifier 2 in terms of both accuracy score and precision score, the classifier is rf\_default.

As for the eli5 charts, the first image below is the lr model and the second one is rf. Comparing the two images, we can find, in lr model, positive features have higher weights, for instance, “excellent”, “perfect” and “funniest” are at the top with highest weights, while negative sentiments like “worst” and “waste” are at the bottom with lowest weights.

However, for the rf model, most of the negative sentiment features have high weights, like “bad” and “worst” are the two tops with highest weights.

```
In [23]: eli5.show_weights(classifier1, top=25)
```

Out[23]: y=1 top features

Weight?	Feature
+0.782	excellent
+0.725	perfect
+0.715	funniest
+0.701	superb
+0.682	refreshing
...	34420 more positive ...
...	33966 more negative ...
-0.668	lame
-0.675	lacks
-0.684	unfortunately
-0.716	worse
-0.722	poor
-0.729	laughable
-0.730	dull
-0.733	disappointing
-0.735	ridiculous
-0.740	save
-0.742	badly
-0.757	avoid
-0.795	mess
-0.844	horrible
-0.881	boring
-0.954	poorly
-1.031	disappointment
-1.087	awful
-1.198	worst
-1.336	waste

```
In [25]: eli5.show_weights(classifier2, top=25)
```

Out[25]:

Weight	Feature
0.0105 ± 0.0297	bad
0.0093 ± 0.0233	worst
0.0071 ± 0.0169	great
0.0049 ± 0.0140	waste
0.0047 ± 0.0141	awful
0.0042 ± 0.0048	and
0.0034 ± 0.0078	no
0.0031 ± 0.0024	the
0.0030 ± 0.0073	nothing
0.0030 ± 0.0083	boring
0.0029 ± 0.0082	terrible
0.0029 ± 0.0075	acting
0.0028 ± 0.0080	excellent
0.0027 ± 0.0025	of
0.0027 ± 0.0022	is
0.0026 ± 0.0069	money
0.0026 ± 0.0057	just
0.0025 ± 0.0073	minutes
0.0025 ± 0.0058	wonderful
0.0025 ± 0.0054	plot
0.0025 ± 0.0044	best
0.0025 ± 0.0026	this
0.0025 ± 0.0027	was
0.0024 ± 0.0077	worse
0.0024 ± 0.0021	in
...	68385 more ...

As for the errors generated in the notebook, cases where lr\_default and rf\_default disagree is 5.008%, cases where rf\_default and lr\_custom disagree is 15.572%, cases where lr\_custom and rf\_custom disagree is 15.774%, and cases where lr\_default and rf\_custom disagree is 5.06%. Generally, there's less errors between the two default classifiers, and also less errors in the random forest classifiers. In terms of the impact of custom features, current custom features did not help a lot in reducing the errors, so future actions could better transform the data and optimize different features to reduce the errors.

#### **IV. Statement of Collaborations**

We followed the academic honesty guidelines posted on the course website. We organized to discuss the task descriptions, requirements, bugs in our code, and the relevant technical content *before* we start working on it.